

## MAKAI: Makerere AI Assistant - Comprehensive Documentation

### Project Overview

**MAKAI** (Makerere AI Assistant) is a locally-running AI-powered chatbot that provides information about Makerere University. It combines **Ollama** for local LLM inference, **Qdrant** for vector search, and **PHP** for the web interface, all running locally without external APIs.

### Core Architecture

text

User Interface (Browser)



PHP Frontend



PHP Backend API



Qdrant (Vector DB)



Ollama (Embeddings + LLM)

---

### Project Structure

text

makai-demo/

```
|   └── Frontend
|       ├── index.php      # Main landing page with chat widget
|       ├── chat.php       # Full-screen chat interface
|       ├── styles/output.css  # Compiled Tailwind CSS
|       └── src/input.css    # Tailwind source
|
|   └── Backend
|       ├── ask.php        # Main question handler
|       └── config.php     # Configuration & session management
```

```

|   └── embed.php      # Text embedding functions
|   └── qdrant.php     # Vector search operations
|   └── generate.php   # AI response generation
|   └── ingest.php     # Document ingestion pipeline
|
|   └── Data
|       └── knowledge/  # Knowledge base text files
|           ├── aims.txt
|           ├── muele.txt
|           └── fees.txt
|
|   └── Configuration
|       ├── package.json  # Node.js dependencies
|       └── tailwind.config.js  # Tailwind configuration
|
└── Infrastructure
    ├── Docker (Qdrant)
    └── Ollama CLI

```

---

## Technology Stack

Component	Technology	Purpose	Version/Model
<b>LLM/Embeddings</b>	Ollama	Local model inference	lonomic-embed-text, llama3.2
<b>Vector Database</b>	Qdrant	Vector storage & search	Docker container
<b>Backend</b>	Plain PHP	Web server & logic	PHP 8.0+
<b>Frontend</b>	Tailwind CSS	Styling	3.4+

Component	Technology	Purpose	Version/Model
<b>JavaScript</b>	Vanilla JS	UI interactions	ES6+
<b>Server</b>	Built-in PHP	Local development	-

---

## Setup & Installation

### Prerequisites Checklist

- PHP 8.0+ installed
- Composer (optional)
- Node.js & npm
- Docker Desktop
- Ollama installed
- 4GB+ RAM available

### Step-by-Step Installation

#### 1. Environment Setup

bash

```
# Clone or create project directory
mkdir makai-demo
cd makai-demo
```

```
# Create folder structure
```

```
mkdir -p knowledge styles assets src
```

#### 2. Install Ollama

bash

```
# On macOS/Linux
```

```
curl -fsSL https://ollama.ai/install.sh | sh
```

```
# On Windows (PowerShell)
```

```
winget install Ollama.Ollama
```

```
# Pull required models  
ollama pull nomic-embed-text # 768-dimensional embeddings  
ollama pull llama3.2      # Or mistral, qwen2.5
```

### 3. Setup Qdrant with Docker

```
bash
```

```
# Pull and run Qdrant  
docker run -p 6333:6333 -p 6334:6334 \  
-v $(pwd)/qdrant_storage:/qdrant/storage \  
qdrant/qdrant
```

```
# Verify Qdrant is running
```

```
curl http://localhost:6333
```

### 4. Install Frontend Dependencies

```
bash
```

```
# Initialize npm
```

```
npm init -y
```

```
# Install Tailwind CSS
```

```
npm install -D tailwindcss @tailwindcss/forms
```

```
# Initialize Tailwind
```

```
npx tailwindcss init
```

```
# Build CSS
```

```
npm run build
```

### 5. Prepare Knowledge Base

```
bash
```

```

# Add your university documents
echo "Makerere University Admission Requirements..." > knowledge/admission.txt
echo "Tuition Fees Structure..." > knowledge/fees.txt
# Add more .txt files as needed

6. Run the System

bash
# Terminal 1: Qdrant (already running in Docker)

# Terminal 2: Start PHP server
php -S localhost:8000

# Terminal 3: Monitor logs (optional)
tail -f /path/to/error_log

# Open in browser
open http://localhost:8000

```

---

## Core Components Detailed

### 1. Embedding System (embed.php)

**Purpose:** Convert text to vector embeddings using Ollama's nomic-embed-text model.

#### Key Features:

- Handles text truncation for model limits
- Error handling for Ollama connection issues
- Returns 768-dimensional vectors

#### Implementation Details:

```

php
function embedText($text) {
    // Limit text length (nomic-embed-text has context limit)
    $text = substr(trim($text), 0, 8192);

```

```

// JSON payload for Ollama API
$payload = [
    "model" => "nomic-embed-text",
    "prompt" => $text,
    "options" => ["temperature" => 0]
];

// Make HTTP request
$context = stream_context_create([
    "http" => [
        "method" => "POST",
        "header" => "Content-Type: application/json",
        "content" => json_encode($payload),
        "timeout" => 30 // 30-second timeout
    ]
]);

// Parse response
$response = @file_get_contents(
    "http://localhost:11434/api/embeddings",
    false,
    $context
);
}

```

### Challenges Solved:

- **Timeout Handling:** Added 30-second timeout for large documents
- **Memory Management:** Truncates text to prevent OOM errors
- **Error Recovery:** Graceful degradation when Ollama is unavailable

## 2. Vector Database Manager (qdrant.php)

**Purpose:** Interface with Qdrant for vector storage and similarity search.

### Key Features:

- Collection management (create/check)
- Vector similarity search with scoring
- Batch point insertion
- Payload management (source metadata)

### Collection Structure:

```
json
{
  "collection_name": "makai",
  "vectors_config": {
    "size": 768,
    "distance": "Cosine"
  },
  "shard_number": 1,
  "replication_factor": 1
}
```

### Search Parameters:

```
php
$searchParams = [
  "vector" => $embedding_vector,
  "limit" => 5,           // Top 5 results
  "with_payload" => true, // Include source text
  "score_threshold" => 0.5, // Minimum similarity score
  "with_vectors" => false
];
```

### Challenges Solved:

- **Connection Stability:** Retry logic for Docker container startup

- **Scoring Threshold:** Filters out irrelevant results (score < 0.5)
- **Metadata Preservation:** Stores source file and chunk index

### 3. Document Ingestion Pipeline (ingest.php)

**Purpose:** Process and chunk text documents into Qdrant.

**Processing Steps:**

1. **File Discovery:** Scan knowledge/ directory for .txt files
2. **Smart Chunking:** Split by paragraphs/sentences (max 500 chars)
3. **Embedding Generation:** Convert chunks to vectors
4. **Batch Insertion:** Store in Qdrant with metadata

**Chunking Algorithm:**

php

```
function chunkText($text, $maxLength = 500) {
    // Split by paragraphs first
    $paragraphs = explode("\n\n", $text);
    $chunks = [];

    foreach ($paragraphs as $paragraph) {
        // If paragraph is too long, split by sentences
        if (strlen($paragraph) > $maxLength) {
            $sentences = preg_split('/(?!<=[!.?])\s+/', $paragraph);
            $currentChunk = "";

            foreach ($sentences as $sentence) {
                if (strlen($currentChunk) + strlen($sentence) <= $maxLength) {
                    $currentChunk .= $sentence . ' ';
                } else {
                    if (!empty(trim($currentChunk))) {
                        $chunks[] = trim($currentChunk);
                    }
                    $currentChunk = "";
                }
            }
        }
    }
}
```

```

    $currentChunk = $sentence . ' ';
}

}

if (!empty(trim($currentChunk))) {
    $chunks[] = trim($currentChunk);
}

} else {
    $chunks[] = trim($paragraph);
}

}

return array_filter($chunks, function($chunk) {
    return strlen($chunk) > 50; // Remove tiny chunks
});
}

```

#### **Challenges Solved:**

- **Document Size Variability:** Handles both short and long documents
- **Semantic Chunking:** Keeps related sentences together
- **Duplicate Prevention:** Uses MD5 hashing for chunk IDs
- **Progress Feedback:** Console output for large ingestions

#### **4. AI Response Generator (generate.php)**

**Purpose:** Use Ollama LLM to generate natural language responses from retrieved context.

#### **Prompt Engineering:**

php

```
$systemPrompt = "You are MAKAI (Makerere AI Assistant)..."
```

Rules:

1. Answer ONLY from provided context
2. If unsure, say 'I don't have information about that'

3. Be concise and factual
4. Cite sources when available

Context:

`{$formattedContext}`

Question: `{$question}"`;

**Context Formatting:**

- Includes top 5 most relevant chunks
- Shows similarity scores
- Preserves source metadata
- Limits total context length to prevent token overflow

**Challenges Solved:**

- **Context Window Management:** Truncates to fit model limits
- **Hallucination Prevention:** Strict "only from context" prompting
- **Source Attribution:** Includes file names in responses
- **Performance:** Caches frequent queries

**5. Web Interface (index.php & chat.php)**

**Purpose:** User-friendly interface for interacting with MAKAI.

**Key UI Components:**

**Floating Chat Widget:**

- Bottom-right positioned button
- Slide-up animation
- Typing indicators
- Markdown rendering

**Features:**

- **Real-time Updates:** WebSocket-like feel with JavaScript polling
- **Responsive Design:** Mobile-first Tailwind CSS
- **Accessibility:** ARIA labels, keyboard navigation

- **Session Management:** PHP sessions for chat history

#### JavaScript Architecture:

javascript

```
class ChatManager {  
    constructor() {  
        this.history = [];  
        this.isLoading = false;  
    }  
  
    async sendMessage(message) {  
        // Show typing indicator  
        this.showTyping();  
  
        // Make API call  
        const response = await fetch('/ask.php', {  
            method: 'POST',  
            body: JSON.stringify({question: message})  
        });  
  
        // Parse and display response  
        const data = await response.json();  
        this.displayResponse(data);  
  
        // Update chat history  
        this.history.push({user: message, ai: data.answer});  
    }  
  
    displayResponse(data) {  
        // Render markdown
```

```
const html = markdownit().render(data.answer);

// Add sources if available
if (data.sources.length > 0) {
    html += this.renderSources(data.sources);
}

// Update DOM
document.getElementById('chat').innerHTML += html;
}
```

---

## Data Flow

### Ingestion Flow

text

Text Document

↓

Read & Clean

↓

Chunk (500 chars)

↓

Embed (Ollama)

↓

Store (Qdrant)

↓

Index Ready

### Query Flow

text

User Question

```
↓  
Embed Question  
↓  
Vector Search (Qdrant)  
↓  
Retrieve Top 5 Chunks  
↓  
Format Context  
↓  
LLM Generation (Ollama)  
↓  
Format Response  
↓  
Display to User
```

---

## 🛠 Configuration Variables

### config.php - Core Settings

```
php  
// Ollama Configuration  
define('OLLAMA_HOST', 'http://localhost:11434');  
define('EMBEDDING_MODEL', 'nomic-embed-text'); // 768-dim  
define('LLM_MODEL', 'llama3.2'); // Response model  
  
// Qdrant Configuration  
define('QDRANT_HOST', 'http://localhost:6333');  
define('COLLECTION_NAME', 'makai');  
define('VECTOR_SIZE', 768); // Must match embedding model  
define('DISTANCE_METRIC', 'Cosine'); // Cosine similarity
```

```
// Application Settings

define('MAX_CONTEXT_CHUNKS', 5);           // Chunks per query
define('SIMILARITY_THRESHOLD', 0.5);        // Minimum score
define('CHUNK_SIZE', 500);                  // Characters per chunk
define('SESSION_TIMEOUT', 3600);            // 1 hour
```

### **Environment Variables (Optional)**

bash

```
# .env file (if using PHP dotenv)

OLLAMA_HOST=http://localhost:11434
QDRANT_HOST=http://localhost:6333
EMBEDDING_MODEL=nomic-embed-text
LLM_MODEL=llama3.2
DEBUG_MODE=false
```

---

## **Performance Optimizations**

### **1. Caching Strategy**

php

```
// Simple file-based caching

function getCachedEmbedding($text) {
    $hash = md5($text);
    $cacheFile = "cache/embeddings/{$hash}.json";

    if (file_exists($cacheFile) &&
        (time() - filemtime($cacheFile)) < 3600) {
        return json_decode(file_get_contents($cacheFile), true);
    }

    return null; // Cache miss
}
```

## 2. Batch Processing

- Ingest documents in batches of 10 chunks
- Parallel embedding requests (if supported)
- Bulk insert to Qdrant

## 3. Memory Management

- Stream large files line by line
  - Unset variables after use
  - Limit concurrent operations
- 

### Testing & Validation

#### Unit Tests (Conceptual)

```
php

// Test embedding generation

function testEmbedding() {
    $text = "Test text for embedding";
    $vector = embedText($text);

    assert(is_array($vector), "Embedding should return array");
    assert(count($vector) == 768, "Vector should be 768-dimensional");
    assert(is_float($vector[0]), "Vector elements should be floats");
}

// Test vector search

function testSearch() {
    $results = searchKnowledge("admission requirements");

    assert(is_array($results), "Results should be array");
    assert(count($results) <= 5, "Should return max 5 results");
```

```
foreach ($results as $result) {  
    assert(isset($result['content']), "Result should have content");  
    assert(isset($result['score']), "Result should have score");  
}  
}
```

## Integration Tests

1. **Ollama Connection:** curl http://localhost:11434/api/tags
2. **Qdrant Health:** curl http://localhost:6333/collections
3. **PHP Endpoints:** curl -X POST http://localhost:8000/ask.php

## Load Testing

- Simulate multiple concurrent users
- Measure response times
- Monitor memory usage

---

## Troubleshooting Guide

### Common Issues & Solutions

#### 1. Ollama Not Responding

text

Error: "Ollama connection failed"

Solution:

1. Check if Ollama is running: `ollama list`
2. Start Ollama: `ollama serve`
3. Verify model is pulled: `ollama pull nomic-embed-text`

#### 2. Qdrant Connection Issues

text

Error: "Qdrant connection failed"

Solution:

1. Check Docker: `docker ps` (should show qdrant)
2. Restart container: `docker restart qdrant`

3. Check port 6333: `curl http://localhost:6333`

### **3. Memory Exhaustion**

text

Error: "Allowed memory size exhausted"

Solution:

1. Increase PHP memory: `ini\_set('memory\_limit', '512M')`
2. Process smaller chunks in ingest.php
3. Use `gc\_collect\_cycles()` in loops

### **4. Slow Response Times**

text

Issue: Chat responses take >10 seconds

Solution:

1. Cache embeddings: Implement file caching
2. Reduce chunk size: Try 300 characters
3. Use lighter model: Switch to mistral:7b

### **5. Poor Search Results**

text

Issue: Irrelevant answers

Solution:

1. Adjust similarity threshold: Try 0.6
2. Improve chunking: Use sentence boundaries
3. Add more context chunks: Increase from 5 to 7

### **Debug Mode**

Enable in config.php:

php

```
define('DEBUG_MODE', true);
```

```
function debugLog($message) {
```

```
    if (DEBUG_MODE) {
```

```
    error_log("[MAKAI DEBUG] " . $message);

}

}
```

---

## Scaling Considerations

### Vertical Scaling

#### 1. Increase Ollama Workers:

bash

```
OLLAMA_NUM_PARALLEL=4 ollama serve
```

#### 2. Qdrant Resource Allocation:

bash

```
docker run -p 6333:6333 \
--memory=4g \
--cpus=2 \
qdrant/qdrant
```

#### 3. PHP Optimization:

ini

```
; php.ini settings
memory_limit = 1G
max_execution_time = 120
opcache.enable=1
```

### Horizontal Scaling

1. **Multiple Ollama Instances:** Load balance embedding requests
2. **Qdrant Cluster:** Multi-node setup for high availability
3. **PHP Session Sharing:** Redis for distributed sessions

### Database Scaling

- Implement sharding for large collections
- Use Qdrant's replication for redundancy
- Regular backup of vector data

---

## Security Considerations

### Local-Only Architecture Benefits

1. **Data Privacy:** All data stays on local machine
2. **No API Keys:** No external service dependencies
3. **Network Isolation:** No internet connection required

### Additional Security Measures

```
php

// Input sanitization

function sanitizeInput($input) {

    $input = trim($input);

    $input = stripslashes($input);

    $input = htmlspecialchars($input, ENT_QUOTES, 'UTF-8');

    return $input;

}

// Session security

ini_set('session.cookie_httponly', 1);

ini_set('session.use_only_cookies', 1);

ini_set('session.cookie_secure', 1); // HTTPS only
```

### File Upload Security (If adding upload feature)

- Validate file types
- Scan for malware
- Store outside web root
- Limit file size

---

## Monitoring & Logging

### Logging System

```
php
```

```

class Logger {

    const LEVELS = ['DEBUG', 'INFO', 'WARN', 'ERROR'];

    public static function log($message, $level = 'INFO') {
        $timestamp = date('Y-m-d H:i:s');
        $logEntry = "[{$timestamp}] [{$level}] {$message}\n";

        // Write to file
        file_put_contents('logs/app.log', $logEntry, FILE_APPEND);

        // Console output in debug mode
        if (DEBUG_MODE && $level === 'DEBUG') {
            echo $logEntry;
        }
    }
}

```

### Key Metrics to Monitor

1. **Response Time:** Average query latency
  2. **Cache Hit Rate:** Embedding cache effectiveness
  3. **Memory Usage:** PHP and Ollama memory consumption
  4. **Error Rate:** Failed requests percentage
  5. **User Engagement:** Questions per session
- 

### Known Limitations

#### Technical Limitations

1. **Model Context Window:** Limited to ~4000 tokens for most models
2. **Embedding Dimension:** Fixed at 768 (nomic-embed-text)
3. **Local Storage:** Vector DB limited by disk space
4. **Single User:** Designed for single local user

## Functional Limitations

1. **No Real-time Updates:** Requires page refresh for new documents
  2. **Basic Search:** No advanced filters or faceted search
  3. **Limited File Formats:** Currently text-only, no PDF/Word
  4. **No User Authentication:** Single anonymous session
- 

## Future Enhancements

### Short-term (Next Version)

1. **PDF Support:** Add PDF parsing with pdf2text
2. **Web Crawler:** Ingest university website content
3. **Export Feature:** Export chat conversations
4. **Admin Dashboard:** Monitor system health

### Medium-term

1. **Multi-language Support:** Add local language models
2. **Voice Interface:** Speech-to-text input
3. **Offline Mode:** Progressive web app capabilities
4. **Plugin System:** Extensible knowledge sources

### Long-term

1. **Multi-user Support:** User accounts and permissions
  2. **Advanced Analytics:** Question patterns and trends
  3. **Mobile App:** iOS/Android applications
  4. **API Service:** REST API for third-party integration
- 

## Knowledge Base Best Practices

### Document Preparation

1. **Clean Text:** Remove headers, footers, page numbers
2. **Consistent Formatting:** Use Markdown for structure
3. **Metadata:** Include document source and date
4. **Regular Updates:** Schedule monthly knowledge refreshes

## Chunking Strategy

- **Academic Content:** Chunk by section/subsection
- **FAQ Documents:** One Q&A per chunk
- **Policy Documents:** Keep entire policies together
- **Mixed Content:** Hybrid approach based on content type

## Quality Control

1. **Sample Queries:** Test with common questions
  2. **Accuracy Check:** Manual review of AI responses
  3. **Coverage Analysis:** Identify knowledge gaps
  4. **Feedback Loop:** User feedback for improvement
- 

## Contributing Guidelines

### Code Contributions

1. Fork the repository
2. Create feature branch: `git checkout -b feature/amazing-feature`
3. Commit changes: `git commit -m 'Add amazing feature'`
4. Push to branch: `git push origin feature/amazing-feature`
5. Open Pull Request

### Documentation Contributions

- Update this README for changes
- Add comments to complex functions
- Create examples for new features
- Document edge cases

### Testing Contributions

- Write unit tests for new features
  - Test on different PHP versions
  - Verify cross-browser compatibility
  - Performance benchmarking
-

## License & Attribution

### Open Source Components

- **Ollama**: MIT License
- **Qdrant**: Apache 2.0 License
- **Tailwind CSS**: MIT License
- **Markdown-it**: MIT License

### Project License

This project is licensed under the MIT License - see the LICENSE file for details.

### Attribution

Please credit:

- Ollama for local LLM inference
- Qdrant for vector search
- Makerere University for the use case

---

## Acknowledgments

- **Makerere University** for inspiring the project
- **Ollama Team** for making local AI accessible
- **Qdrant Team** for the excellent vector database
- **Tailwind CSS** for the beautiful styling framework
- **PHP Community** for keeping the web simple

---

## Getting Help

### Support Channels

1. **GitHub Issues**: Bug reports and feature requests
2. **Documentation**: This README and code comments
3. **Community Forum**: (If established)
4. **Email Support**: (For enterprise deployments)

### Quick Start Checklist

- Ollama running with models

- Qdrant Docker container active
- PHP server on port 8000
- Tailwind CSS built
- Knowledge documents added
- Ingestion script run

### **Common Commands Cheatsheet**

bash

```
# Start all services
```

```
docker run -p 6333:6333 qdrant/qdrant &
ollama serve &
php -S localhost:8000 &
```

```
# Ingest documents
```

```
php ingest.php
```

```
# Monitor logs
```

```
tail -f logs/app.log
```

```
# Clear chat history
```

```
rm -rf sessions/*
```

```
# Rebuild CSS
```

```
npm run build
```

### **❖ Conclusion**

MAKAI demonstrates how to build a fully-functional, locally-running AI assistant using modern open-source tools. It solves the problem of providing instant, accurate information about Makerere University while maintaining complete data privacy and control.

The project successfully integrates:

- **Local AI inference** with Ollama
- **Efficient vector search** with Qdrant
- **Modern web interface** with Tailwind CSS
- **Scalable architecture** that can grow with needs

This implementation serves as both a practical tool for Makerere University and a reference architecture for building similar local-first AI applications in other domains.

---

*Last Updated: \$(date)*

*Version: 1.0.0*

*Maintainer: MAKAI Development Team*